

## Results

No defaults (# of epochs, batch size, etc.) were changed from original code.

All optimizers from <https://pytorch.org/docs/stable/optim.html#algorithms>.

Optimizer	Accuracy	Loss	Recall	F1	AUC
Adam	.718	.639	.706	.612	.772
Adamax	.698	.643	.706	.603	.770
RMSprop	.769	.632	.687	.636	.766
SGD w/ Momentum	.638	.691	.606	.532	.649
Adadelta	.145	.693	.500	.127	.482
Adagrad	.690	.664	.678	.587	.743
NAdam	.770	.643	.689	.637	.770
RAdam	.650	.656	.698	.571	.766
Rprop	.692	.600	.706	.599	.771
ASGD	.540	.693	.559	.467	.582

Adam = Default optimizer used in paper

NAdam = Biggest improvement

### Adam vs NAdam Results

Optimizer	Accuracy	Loss	Recall	F1	AUC
Adam	.718	.639	.706	.612	.772
NAdam	.770	.643	.689	.637	.770
Difference	+6.99%	+6.2%	-2.44%	+4%	-.26%

### Summary:

With NAdam improving the accuracy by almost 7% and F1 score increasing by 4%, it was enough to show a considerable improvement, despite the minor increase in loss (.62%) and minor decreases in recall (2.4%) and AUC (.26%).